

Multi-Object Synchronization

Synchronization Performance

- A program with lots of concurrent threads can still have poor performance on a multiprocessor:
 - Overhead of creating threads, if not needed
 - Lock contention: only one thread at a time can hold a given lock
 - Shared data protected by a lock may ping back and forth between cores
 - False sharing: communication between cores even for data that is not shared

Atomic CompareAndSwap

- Operates on a memory word
- Check that the value of the memory word hasn't changed from what you expect
 - E.g., no other thread did compareAndSwap first
- If it has changed, return an error (and loop)
- If it has not changed, set the memory word to a new value

Multi-Object Automicity

- Combining into one object, e.g., using a function
 - `checkforMilkAndSetNoteIfNeeded()`
- Acquire-All/Release All (Serializable)
 - AND SYNCHRONIZATION
 - Must know all required locks *a priori*
- Two-Phase Locking
 - Acquire locks as needed and do not release locks until all locks have been acquired
 - How to avoid deadlock?

Deadlock Definition

- Resource: any (passive) thing needed by a thread to do its job (CPU, disk space, memory, lock)
 - Preemptable: can be taken away by OS
 - Non-preemptable: must leave with thread
- Starvation: thread waits indefinitely
- Deadlock: circular waiting for resources
 - Deadlock => starvation, but not vice versa

Example: two locks

Thread A

```
lock1.acquire();  
lock2.acquire();  
lock2.release();  
lock1.release();
```

Thread B

```
lock2.acquire();  
lock1.acquire();  
lock1.release();  
lock2.release();
```

Bidirectional Bounded Buffer

Thread A

```
buffer1.put(data);  
buffer1.put(data);
```

```
buffer2.get();  
buffer2.get();
```

Thread B

```
buffer2.put(data);  
buffer2.put(data);
```

```
buffer1.get();  
buffer1.get();
```

Suppose buffer1 and buffer2 both start almost full.

Two locks and a condition variable


Thread A

```
lock1.acquire();  
...  
lock2.acquire();  
while (need to wait) {  
    condition.wait(lock2);  
}  
lock2.release();  
...  
lock1.release();
```


Thread B

```
lock1.acquire();  
...  
lock2.acquire();  
...  
condition.signal(lock2);  
...  
lock2.release();  
...  
lock1.release();
```


Yet another Example



Dining Lawyers



Each lawyer needs two chopsticks to eat.
Each grabs chopstick on the right first.

Necessary Conditions for Deadlock

- Limited access to resources
 - If infinite resources, no deadlock!
- No preemption
 - If resources are virtual, can break deadlock
- Multiple independent requests
 - “wait while holding”
- Circular chain of requests

Question

- How does Dining Lawyers meet the necessary conditions for deadlock?
 - Limited access to resources
 - No preemption
 - Multiple independent requests (wait while holding)
 - Circular chain of requests
- How can we modify Dining Lawyers to prevent deadlock?

Preventing Deadlock

- Exploit or limit program behavior
 - Limit program from doing anything that might lead to deadlock
- Predict the future
 - If we know what program will do, we can tell if granting a resource might lead to deadlock
- Detect and recover
 - If we can rollback a thread, we can fix a deadlock once it occurs

Exploit or Limit Behavior

- Provide enough resources
 - How many chopsticks are enough?
- Eliminate wait while holding
 - Release lock when calling out of module
 - Telephone circuit setup
- Eliminate circular waiting
 - Lock ordering: always acquire locks in a fixed order
 - Example: move file from one directory to another

Example

Thread 1

1. Acquire A
- 2.
3. Acquire C
- 4.
5. If (maybe) Wait for B

Thread 2


- 1.
2. Acquire B
- 3.
4. Wait for A

How can we make sure to avoid deadlock?

Deadlock Dynamics

- Safe state:
 - For any possible sequence of future resource requests, it is possible to eventually grant all requests
 - May require waiting even when resources are available!
- Unsafe state:
 - Some sequence of resource requests can result in deadlock
- Doomed state:
 - All possible computations lead to deadlock

Possible System States



Question

- What are the doomed states for Dining Lawyers?
- What are the unsafe states?
- What are the safe states?

Communal Dining Lawyers

- n chopsticks in middle of table
- n lawyers, each can take one chopstick at a time
- What are the safe states?
- What are the unsafe states?
- What are the doomed states?

Communal Mutant Dining Lawyers

- N chopsticks in the middle of the table
- N lawyers, each takes one chopstick at a time
- Lawyers need k chopsticks to eat, $k > 1$

- What are the safe states?
- What are the unsafe states?
- What are the doomed states?

Communal Mutant Absent-Minded Dining Lawyers

- N chopsticks in the middle of the table
- N lawyers, each takes one chopstick at a time
- Lawyers need k chopsticks to eat, $k > 1$
 - k larger if lawyer is talking on his/her cellphone
- What are the safe states?
- What are the unsafe states?
- What are the doomed states?

Predict the Future

- Banker's algorithm
 - State maximum resource needs in advance
 - Allocate resources dynamically when resource is needed -- wait if granting request would lead to deadlock
 - Request can be granted if some sequential ordering of threads is deadlock free

Banker's Algorithm

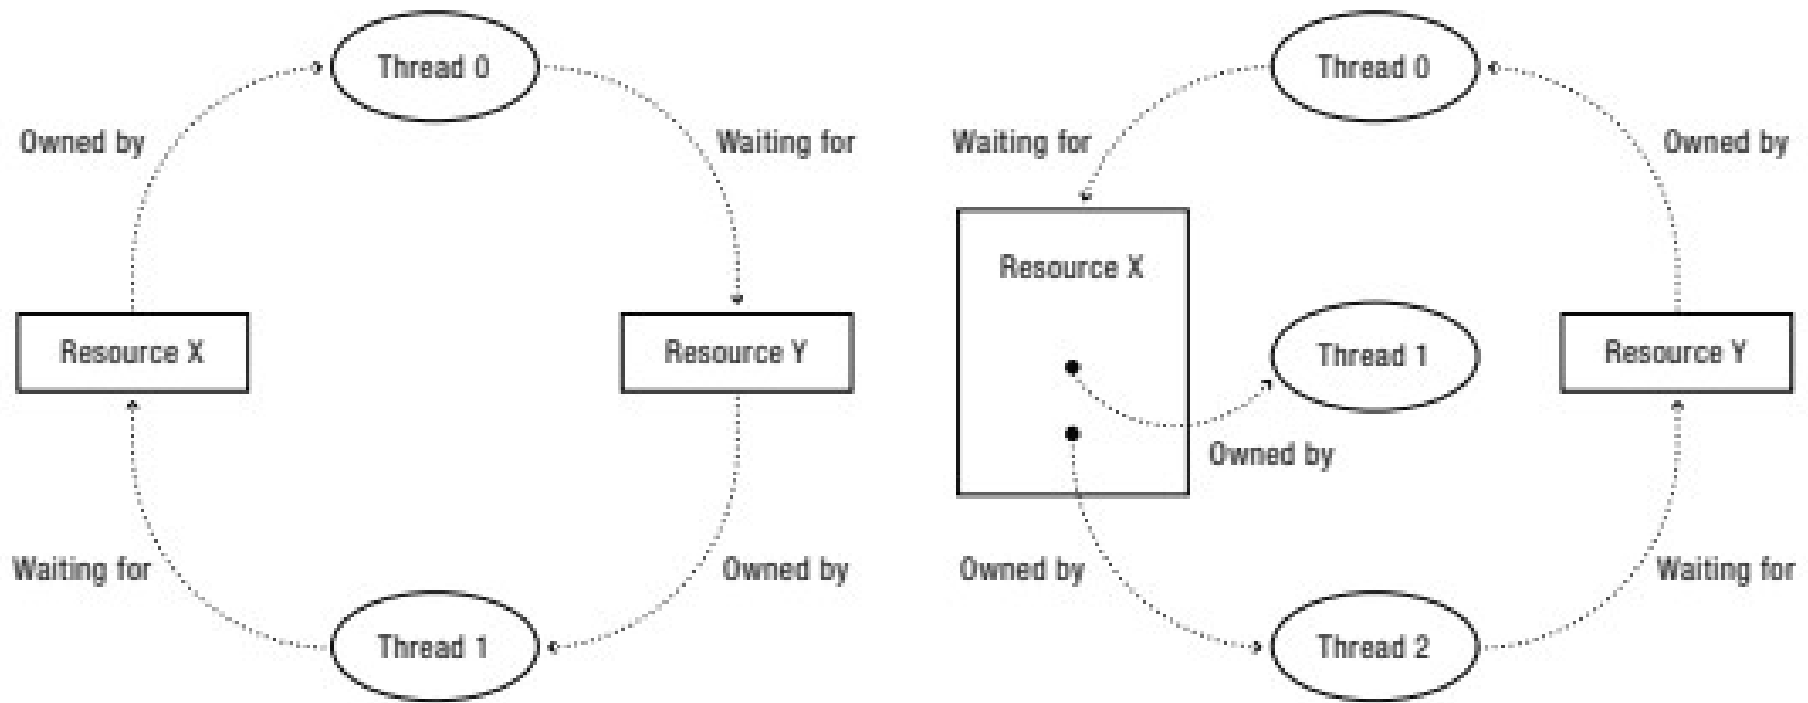
- Grant request iff result is a safe state
- Sum of maximum resource needs of current threads can be greater than the total resources
 - Provided there is some way for all the threads to finish without getting into deadlock
- Example: proceed iff
 - total available resources - # allocated \geq max remaining that might be needed by this thread in order to finish
 - Guarantees this thread can finish

Check Figure 6.20 and 6.21 for detailed pseudo-code.

Detect and Repair

- Algorithm
 - Scan wait for graph
 - Detect cycles
 - Fix cycles
- Proceed without the resource
 - Requires robust exception handling code
- Roll back and retry
 - Transaction: all operations are provisional until have all required resources to complete operation

Detecting Deadlock



An example of a non-deadlocked scenario is mis-classified as being deadlocked

More precise: Figure 6.23

Non-Blocking Synchronization

- Goal: data structures that can be read/modified without acquiring a lock
 - No lock contention!
 - No deadlock!
- General method using compareAndSwap
 - Create copy of data structure
 - Modify copy
 - Swap in new version iff no one else has
 - Restart if pointer has changed

Lock-Free Bounded Buffer

```
tryget() {  
    do {  
        copy = ConsistentCopy(p);  
        if (copy->front == copy->tail)  
            return NULL;  
        else {  
            item = copy->buf[copy->front % MAX];  
            copy->front++;  
        } while (compareAndSwap(&p, p, copy));  
    } while (true);  
    return item;  
}
```